

Daniel Lu

☎ +1 (613) 720-2860 | ✉ daniel.lu@gmail.com | 🌐 daniel.lu.ca | 🐙 finityfly | 📄 daniel-jmlu

Education

Carleton University, B.Sc. in Computer Science (Honours, Co-op)

Ottawa, Canada

AI & Machine Learning Stream - GPA: 3.9/4.0

Sep 2023 – Apr 2028

- **Coursework:** Supervised/Unsupervised Methods, Transformer Architectures, Language Modeling & Embeddings

Experience

Software Developer

May 2026 – Aug 2026

Shopify

Toronto, Canada

- Incoming Summer 2026

Software Developer

May 2025 – Aug 2025

TrendAI

Ottawa, Canada

- Pioneered TrendAI's **first** MCP-based digital twin prototype, architecting a Neo4j graph RAG pipeline for predictive infrastructure security analytics on Vision One, directly securing C-suite approval for full-scale development.
- Improved system latency by **30%** and reduced monthly cloud costs by **18%** across **120+ active enterprise users** via AWS cloud infrastructure optimizations and Kubernetes Horizontal Pod Autoscaling.
- Engineered CI/CD pipelines for internal SDKs and database deployment systems serving **200+ developers** and **10,000+ stakeholders**, reducing release cycle duration by **65%** (**10 hours** → **under 4 hours**).

Software Developer

Jan 2025 – Apr 2025

Blackberry QNX

Ottawa, Canada

- Engineered core automation libraries in Python for the QNX RTOS, expanding total code coverage by **15%** (60% → 75%).
- Reduced system validation cycle time by **19%** by engineering automated testing scripts in Python, C, and Bash for QNX RTOS build pipelines, cutting manual verification effort from **8 hours** to **6.5 hours** per release.

Projects

Mime | [GitHub](#)

- Bridged the visual-audio immersion gap in cross-lingual video calls by engineering a lightweight real-time 3D facial overlay that renders lip-synced translated speech with **91% accuracy**.
- Architected and trained a CNN-Transformer audio encoder on the 76-hour BEAT motion-capture dataset, regressing 52 ARKit facial blendshape coefficients at **0.0098 MSE**.
- Shipped a zero-client-modification pipeline compatible with all major video conferencing platforms, achieving **<250ms** end-to-end streaming latency and **100ms TTFT**.

Tobio | [Website](#) | [GitHub](#)

- Engineered a sensor-free 3D player tracking system from standard game footage using a monocular 3D reconstruction pipeline, deploying the product to **3 university volleyball teams** with no specialized hardware.
- Cut post-game analysis time by **85%** by piping a sliding-window action classifier into a vector-indexed RAG pipeline, enabling natural language querying of plays, rotations, and set patterns.
- Achieved **sub-centimeter tracking precision** on player and ball trajectories by pairing YOLOv11 detection with a ReID module to preserve player identity through occlusion, enabling continuous real-time 3D mapping across full match footage.

Technical Skills

Languages: Python, C/C++, Java, JavaScript, TypeScript, Ruby, SQL

Libraries and Frameworks: Hugging Face, PyTorch, NumPy, Scikit-learn, React, FastAPI

Platforms & Technologies: AWS, Docker, Kubernetes, Jenkins, Terraform, REST APIs, Git